**1.Problem Statement:**

The client wants to predict insurance charges based on several parameters such as age, gender, BMI, number of children, smoking habits. I have a dataset containing these variables, along with the corresponding insurance charges. As a data scientist, my task is to develop a machine learning model that accurately predicts the insurance charges for new data inputs based on these parameters.

**2.Basic info of dataset :**

Here there are 6 columns like age, gender, bmi, children, smoker, charges.

Here there are 1338 rows.

Here we give input as independent=dataset[['age','bmi', 'children','sex_male', 'smoker_yes']]

Here we give output as dependent=dataset[["charges"]]

**4.According to my work done on various algorithm, I got best performance on random forest with 0.8710 r_score.**

Based on the performance scores , the **best model** for this dataset is:

**Random Forest**

- **Criterion**: squared_error
- **Max features**: 'sqrt'
- **n_estimators**: 100
- **Performance Score**: 0.8710

This configuration gives the highest score and the best model compared with other algorithm.

**5.To find machine learning regression method using in r2 value.**

1. **MULTIPLE LINEAR REGRESSION: r_score : 0.7894790349867009**

2. **SUPPORT VECTOR MACHINE:**

| S.NO | HYPER PARAMETER | LINEAR (r_score) | RBF (r_score) | POLY (r_score) | SIGMOID (r_score) |
|---|---|---|---|---|---|
| 1. | C10 | -0.0403 | -0.09506 | -0.1202 | -0.0992 |
| 2. | C100 | 0.5218 | -0.09506 | -0.1319 | -0.1241 |
| 3. | C500 | 0.6133 | -0.15659 | -0.1166 | -0.4197 |
| 4. | C1000 | 0.6188 | -0.14956 | -0.0923 | -1.5217 |
| 5. | C2000 | 0.6253 | -0.14219 | -0.0423 | -5.0920 |
| 6. | C3000 | 0.6662 | -0.13060 | 0.0062 | -10.947 |

3. **DECISION TREE:**

| S.NO | CRITERION | MAX FEATURES | SPLITTER | R_VALUE |
|---|---|---|---|---|
| 1 | squared_error | Sqrt | Random | 0.7313 |
| 2 | squared_error | sqrt | Best | 0.6384 |
| 3 | squared_error | log2 | Best | 0.5672 |
| 4 | squared_error | log2 | Random | 0.6583 |
| 5 | friedman_mse | log2 | Random | 0.6913 |
| 6 | friedman_mse | log2 | Best | 0.7112 |
| 7 | friedman_mse | Sqrt | Best | 0.7152 |
| 8 | friedman_mse | Sqrt | Random | 0.6126 |
| 9 | absolute_error | Sqrt | Random | 0.7426 |
| 10 | absolute_error | Sqrt | Best | 0.7384 |
| 11 | absolute_error | log2 | Best | 0.7609 |
| 12 | absolute_error | log2 | Random | 0.7389 |
| 13 | Poisson | log2 | Random | 0.6561 |
| 14 | Poisson | log2 | Best | 0.7337 |
| 15 | Poisson | Sqrt | Best | 0.7337 |
| 16 | Poisson | sqrt | Random | 0.6697 |

**4. RANDOM FOREST:**

| S.NO | CRITERION | MAX FEATURES | N_ESTIMATORS | R_VALUE |
|---|---|---|---|---|
| **1** | absolute_error | sqrt | 10 | 0.8574 |
| **2** | absolute_error | Log2 | 10 | 0.8574 |
| **3** | absolute_error | Sqrt | 100 | 0.8710 |
| **4** | absolute_error | Log2 | 100 | 0.8710 |
| **5** | friedman_mse | Sqrt | 10 | 0.8502 |
| **6** | friedman_mse | Sqrt | 100 | 0.8710 |
| **7** | friedman_mse | Log2 | 10 | 0.8502 |
| **8** | friedman_mse | Log2 | 100 | 0.8710 |
| **9** | Poisson | Sqrt | 10 | 0.8544 |
| **10** | Poisson | Sqrt | 100 | 0.8680 |
| **11** | Poisson | Log2 | 10 | 0.8544 |
| **12** | Poisson | Log2 | 100 | 0.8680 |
| **13** | squared_error | Sqrt | 10 | 0.8520 |
| **14** | squared_error | Sqrt | 100 | 0.8710 |
| **15** | squared_error | Log2 | 10 | 0.8520 |
| **16** | squared_error | Log2 | 100 | 0.8710 |

The configurations with the highest performance (score: 0.8710) are:

1. absolute_error with sqrt, 100 estimators
2. absolute_error with log2, 100 estimators
3. friedman_mse with sqrt, 100 estimators
4. friedman_mse with log2, 100 estimators
5. squared_error with sqrt, 100 estimators
6. squared_error with log2, 100 estimators

All these configurations give the same highest performance of 0.8710.

**6. FROM THE FOUR MODEL RANDOM FOREST GIVES US THE BEST R_SCORE COMPARED TO OTHERS.**

Random Forest is the best model because:

1. It has the highest score (0.8710).

2. It prevents overfitting by averaging many decision trees.

3. It works well with both simple and complex data.

In short, it's accurate, reliable, and works for many situations.

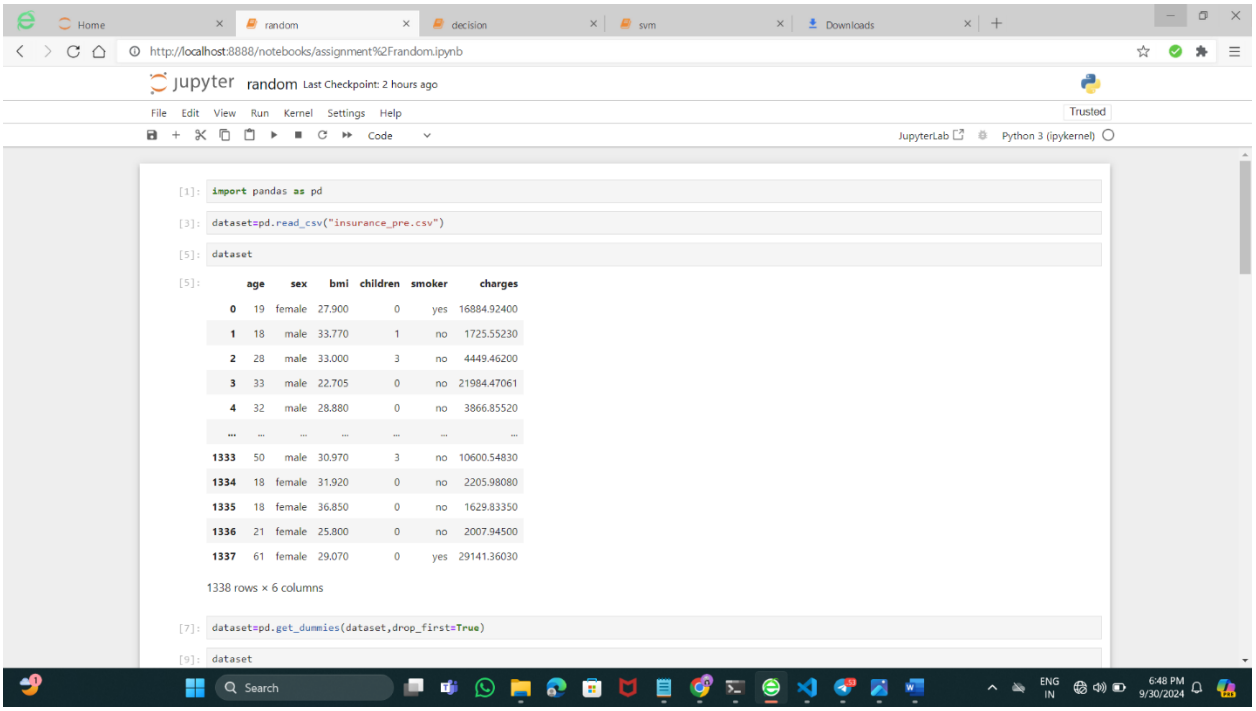
Home
random
decision
svm
Downloads
http://localhost:8888/notebooks/assignment%2Frandom.ipynb
jupyter random Last Checkpoint: 2 hours ago
File Edit View Run Kernel Settings Help
Trusted
Code
JupyterLab
Python 3 (ipykernel)
[1]: import pandas as pd
[3]: dataset=pd.read_csv("insurance_pre.csv")
[5]: dataset
[5]:
age sex bmi children smoker charges
0 19 female 27.900 0 yes 16884.92400
1 18 male 33.770 1 no 1725.55230
2 28 male 33.000 3 no 4449.46200
3 33 male 22.705 0 no 21984.47061
4 32 male 28.880 0 no 3866.85520
... ... ... ... ... ... ...
1333 50 male 30.970 3 no 10600.54830
1334 18 female 31.920 0 no 2205.98080
1335 18 female 36.850 0 no 1629.83350
1336 21 female 25.800 0 no 2007.94500
1337 61 female 29.070 0 yes 29141.36030
1338 rows × 6 columns
[7]: dataset=pd.get_dummies(dataset,drop_first=True)
[9]: dataset
Search
ENG IN
6:48 PM 9/30/2024

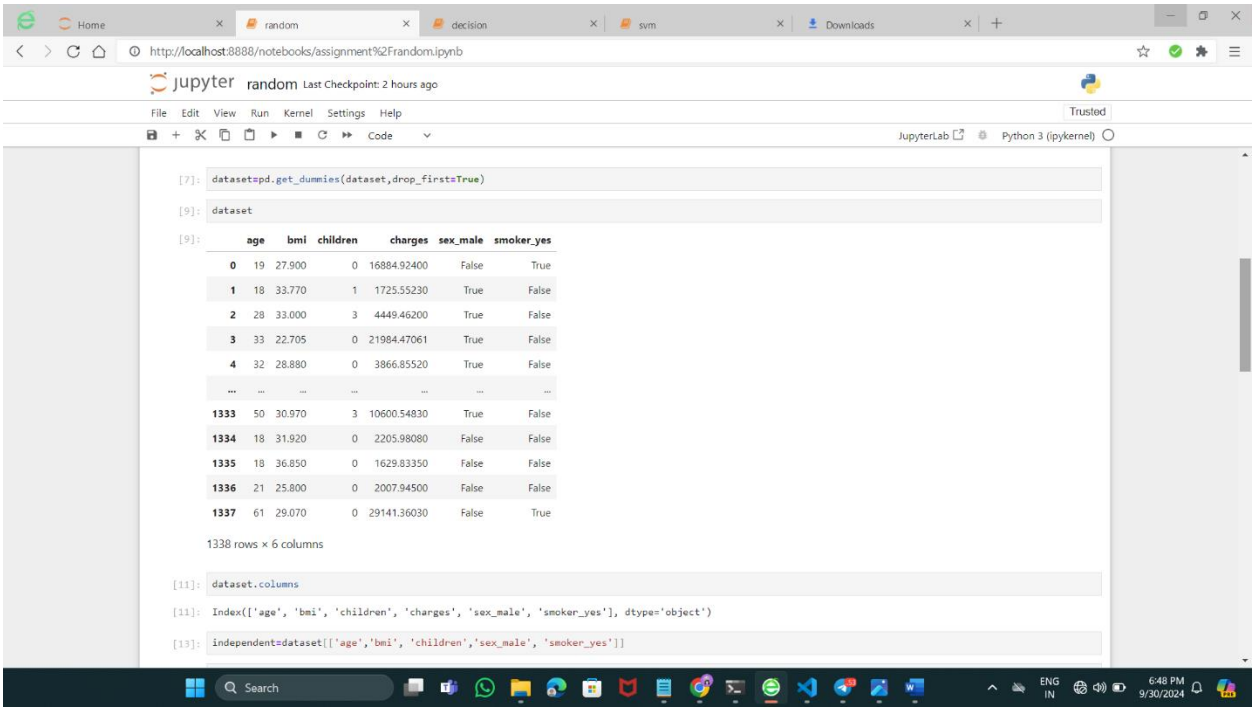
Home
random
decision
svm
Downloads
http://localhost:8888/notebooks/assignment%2Frandom.ipynb
jupyter random Last Checkpoint: 2 hours ago
File Edit View Run Kernel Settings Help
Trusted
Code
JupyterLab
Python 3 (ipykernel)
[7]: dataset=pd.get_dummies(dataset,drop_first=True)
[9]: dataset
[9]:
age bmi children charges sex_male smoker_yes
0 19 27.900 0 16884.92400 False True
1 18 33.770 1 1725.55230 True False
2 28 33.000 3 4449.46200 True False
3 33 22.705 0 21984.47061 True False
4 32 28.880 0 3866.85520 True False
... ... ... ... ... ... ...
1333 50 30.970 3 10600.54830 True False
1334 18 31.920 0 2205.98080 False False
1335 18 36.850 0 1629.83350 False False
1336 21 25.800 0 2007.94500 False False
1337 61 29.070 0 29141.36030 False True
1338 rows × 6 columns
[11]: dataset.columns
[11]: Index(['age', 'bmi', 'children', 'charges', 'sex_male', 'smoker_yes'], dtype='object')
[13]: independent=dataset[['age','bmi', 'children','sex_male', 'smoker_yes']]
Search
ENG IN
6:48 PM 9/30/2024

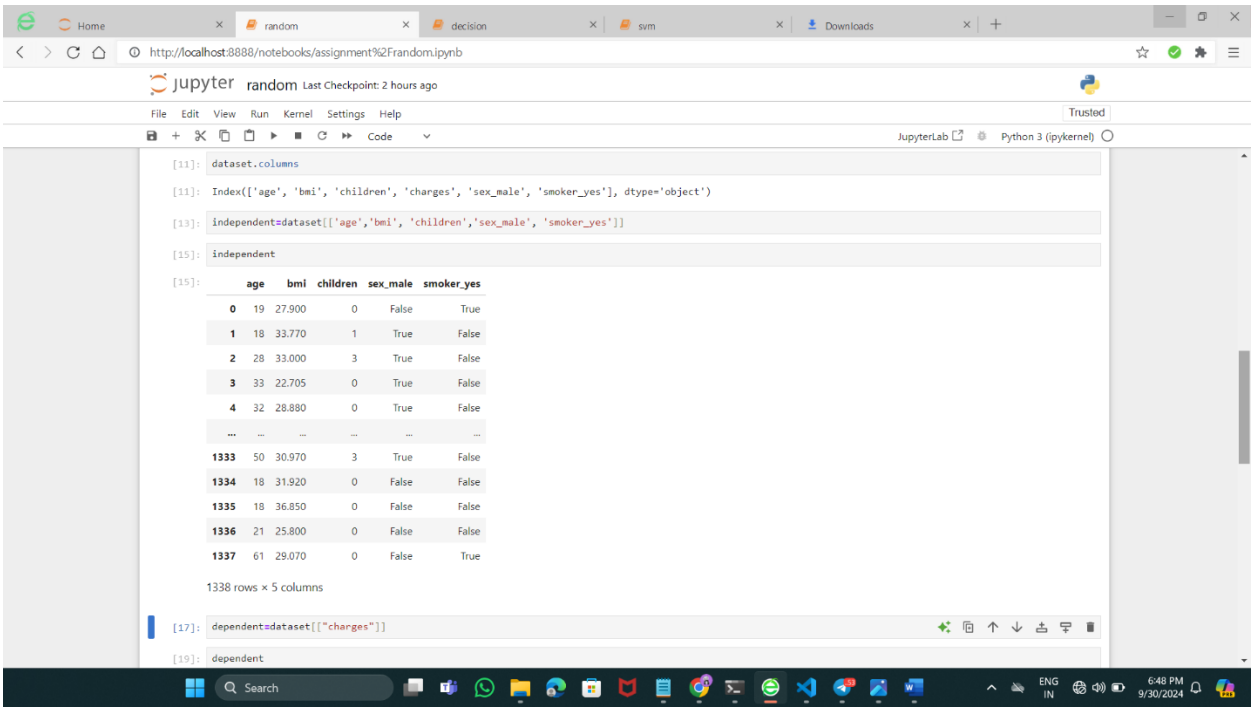
Home
random
decision
svm
Downloads
http://localhost:8888/notebooks/assignment%2Frandom.ipynb
jupyter random Last Checkpoint: 2 hours ago
File Edit View Run Kernel Settings Help
Trusted
Code
JupyterLab
Python 3 (ipykernel)
[11]: dataset.columns
[11]: Index(['age', 'bmi', 'children', 'charges', 'sex_male', 'smoker_yes'], dtype='object')
[13]: independent=dataset[['age','bmi', 'children','sex_male', 'smoker_yes']]
[15]: independent
[15]:
age bmi children sex_male smoker_yes
0 19 27.900 0 False True
1 18 33.770 1 True False
2 28 33.000 3 True False
3 33 22.705 0 True False
4 32 28.880 0 True False
... ... ... ... ... ...
1333 50 30.970 3 True False
1334 18 31.920 0 False False
1335 18 36.850 0 False False
1336 21 25.800 0 False False
1337 61 29.070 0 False True
1338 rows × 5 columns
[17]: dependent=dataset[["charges"]]
[19]: dependent
Search
ENG IN
6:48 PM 9/30/2024

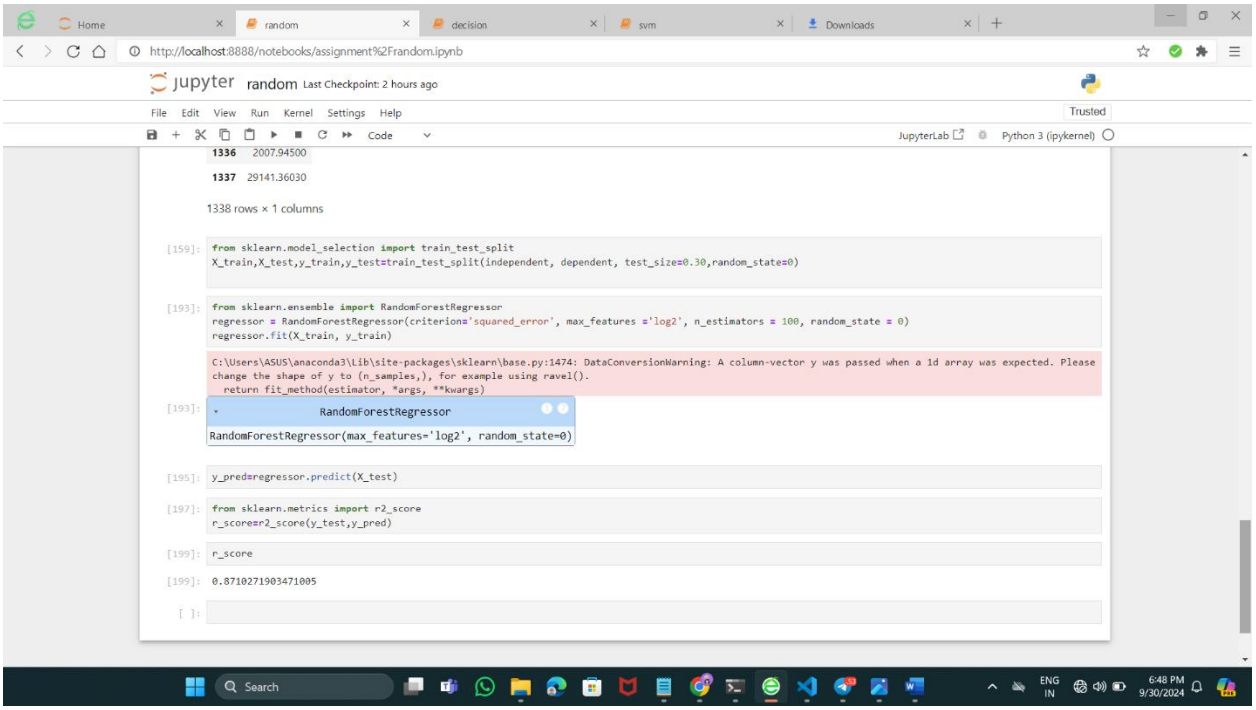
Home
random
decision
svm
Downloads
http://localhost:8888/notebooks/assignment%2Frandom.ipynb
jupyter random Last Checkpoint: 2 hours ago
File Edit View Run Kernel Settings Help
Trusted
Code
JupyterLab
Python 3 (ipykernel)
1336 2007.94500
1337 29141.36030
1338 rows × 1 columns
[159]: from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test=train_test_split(independent, dependent, test_size=0.30,random_state=0)
[193]: from sklearn.ensemble import RandomForestRegressor
regressor = RandomForestRegressor(criterion='squared_error', max_features ='log2', n_estimators = 100, random_state = 0)
regressor.fit(X_train, y_train)
C:\Users\ASUS\anaconda3\Lib\site-packages\sklearn\base.py:1474: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
return fit_method(estimator, *args, **kwargs)
[193]: RandomForestRegressor
RandomForestRegressor(max_features='log2', random_state=0)
[195]: y_pred=regressor.predict(X_test)
[197]: from sklearn.metrics import r2_score
r_score=r2_score(y_test,y_pred)
[199]: r_score
[199]: 0.8710271903471005
[ ]:
Search
ENG IN
6:48 PM 9/30/2024